

STAT 309: MATHEMATICAL COMPUTATIONS I
FALL 2015
LECTURE 1

1. LAUNDRY LIST

- web site: <http://www.stat.uchicago.edu/~lekheng/courses/309/>
- notes: <http://www.stat.uchicago.edu/~lekheng/courses/309/notes/>
- books: <http://www.stat.uchicago.edu/~lekheng/courses/309/books/>
- last year's notes: <http://www.stat.uchicago.edu/~lekheng/courses/309f14/notes/>
- no required textbook
- main references:
 - Trefethen and Bau
 - Watkins
 - Golub and Van Loan
 - Demmel
- facts about matrices:
 - Bernstein
- lectures:
 - no lecture on Tue, Oct 6
 - make-up lectures on Mon, Oct 12 and Mon, Oct 19 in Eckhart 133, 5:30–8:30pm
- homework:
 - homework due beginning of class
 - collaboration allowed but must be declared
 - six assignments, lowest score will be dropped, accounting for 50% of grade
 - no late homework will be accepted
- exams:
 - two quizzes on Thu, Oct 29 and Tue, Nov 24 in Eckhart 133, 3:00–4:20pm
 - in-class, closed-book, no cheat sheet, 2 hr 50 min
- grade: 50% homework, 50% exams
- office hours: Wed, 1:30pm–3:00pm, Eckhart 122

2. NUMERICAL ANALYSIS

- numerical analysis: study of algorithms for continuous mathematics
- examples:
 - linear partial differential equation: given c_α 's, find f

$$\sum_{|\alpha| \leq n} c_\alpha(\mathbf{t}) \frac{\partial^\alpha}{\partial \mathbf{t}^\alpha} f(\mathbf{t}) = 0 \quad (2.1)$$

- Fredholm integral equation of the first kind: given K and g , find f

$$\int_{\Omega} K(\mathbf{s}, \mathbf{t}) f(\mathbf{t}) d\mathbf{t} = g(\mathbf{s}) \quad (2.2)$$

- linear eigenvalue problem: given c_α 's, find f and λ

$$\sum_{|\alpha| \leq n} c_\alpha(\mathbf{t}) \frac{\partial^\alpha}{\partial \mathbf{t}^\alpha} f(\mathbf{t}) = \lambda f(\mathbf{t}) \quad (2.3)$$

- Fredholm integral equation of the first kind: given K and g , find f and λ

$$g(\mathbf{s}) + \lambda \int_{\Omega} K(\mathbf{s}, \mathbf{t}) f(\mathbf{t}) d\mathbf{t} = f(\mathbf{s}) \quad (2.4)$$

- nonlinear optimization: given f_0, \dots, f_m , find \mathbf{t}_{\min}

$$\min f_0(\mathbf{t}) \quad \text{subject to} \quad f_1(\mathbf{t}) \leq 0, \dots, f_m(\mathbf{t}) \leq 0 \quad (2.5)$$

- many scientific and engineering can be formulated in one of these forms — the PDE or integral equations would be a mathematical formulation of physical principles like Newton's second law or Maxwell equations or Schrödinger equation
- we can rarely solve these analytically, i.e., give a useful closed-form formula for the solution
- have to rely on computers, which can only deal with discrete problems
- discretization of (2.1) or (2.2), or Newton method applied to (2.5) yields

$$A\mathbf{x} = \mathbf{b} \quad (2.6)$$

- discretization of (2.3) or (2.4) yields

$$A\mathbf{x} = \lambda \mathbf{x} \quad (2.7)$$

- when we discretize, we have

$$\mathbf{x} = \begin{bmatrix} f(\mathbf{t}_1) \\ f(\mathbf{t}_2) \\ \vdots \\ f(\mathbf{t}_n) \end{bmatrix}$$

- solving for \mathbf{x} gives us a sample of point values of f , which is often enough for many purposes
- the larger n is, the more information we get about f
- the matrix A comes from discretization of the linear operator — differential operators in the case of (2.1) or (2.3) and integral operators in the case of (2.2) or (2.4)
- example: discretizing a 1-dimensional differential operator

$$\frac{d^2}{dt^2} \xrightarrow{\text{discretize}} \begin{bmatrix} 2 & -1 & & \\ -1 & 2 & \ddots & \\ & \ddots & \ddots & -1 \\ & & -1 & 2 \end{bmatrix} \in \mathbb{R}^{n \times n}$$

- example: discretizing a 2-dimensional differential operator

$$\frac{\partial^2}{\partial t_1^2} + \frac{\partial^2}{\partial t_2^2} \xrightarrow{\text{discretize}} \begin{bmatrix} D & -I & & \\ -I & D & \ddots & \\ & \ddots & \ddots & -I \\ & & -I & D \end{bmatrix} \in \mathbb{R}^{mn \times mn}, \text{ where } D = \begin{bmatrix} 4 & -1 & & \\ -1 & 4 & \ddots & \\ & \ddots & \ddots & -1 \\ & & -1 & 4 \end{bmatrix} \in \mathbb{R}^{m \times m}$$

- bottom line: many problems in science and engineering require that we solve (2.6) or (2.7)

3. OPTIMIZATION

- suppose you want to solve an optimization problem

$$\begin{array}{ll} \text{minimize} & f(\mathbf{x}) \\ \text{subject to} & h_i(\mathbf{x}) \leq 0, \quad i = 1, \dots, p, \\ & A\mathbf{x} = \mathbf{b} \end{array}$$

- one of the most widely used algorithm is interior point method (essentially Newton's method adapted to a constrained optimization problem) which requires us to solve a linear system of the form

$$\begin{bmatrix} t\nabla^2 f(\mathbf{x}_k) + \nabla^2 \varphi(\mathbf{x}_k) & A^\top \\ A & 0 \end{bmatrix} \begin{bmatrix} \Delta \mathbf{x}_k \\ \boldsymbol{\nu}_k \end{bmatrix} = - \begin{bmatrix} t\nabla f(\mathbf{x}_k) + \nabla \varphi(\mathbf{x}_k) \\ 0 \end{bmatrix}$$

where φ is the so-called log barrier function that 'traps' the iterates \mathbf{x}_k within the region defined by the constraints

- at each iterate \mathbf{x}_k , we will have to solve such a linear system for $\Delta \mathbf{x}_k$ to obtain the next iterate $\mathbf{x}_{k+1} = \mathbf{x}_k + \Delta \mathbf{x}_k$
- so the computational cost of interior point methods is largely dominated by the cost of solving linear systems

4. MACHINE LEARNING

- many modern problems are information theoretic in nature
 - no differential or integral equations describing your solution f
 - but a large *test set* of given data $\{(x_i, f(x_i)) : i = 1, \dots, n\}$ that allows you to guess your f
- example: classification problems
 - spam identification

$$f : \text{emails} \rightarrow \{\text{spam}, \text{nospam}\}$$

- image recognition

$$f : \text{facial images} \rightarrow \{\text{male}, \text{female}\}$$

or more generally

$$f : \text{handwritten digits} \rightarrow \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$$

- there is no 'Newton's law' type of rule to describe f
- example: supervised learning for binary classification

$$f : X \rightarrow \{-1, +1\}$$

- given training set $\Omega = \{x_1, \dots, x_n\} \subseteq X$, i.e., we already know the value $f(x_i) = y_i$ for any $x_i \in \Omega$
- want to find f , i.e., given some $x \notin \Omega$, we want to predict the value $f(x)$
- let us use spam identification as an example, then for any e-mail $x \in X$,

$$f(x) = \begin{cases} -1 & \text{if } x \text{ is spam} \\ +1 & \text{if } x \text{ is not spam} \end{cases}$$

- we can encode an e-mail as a vector in \mathbb{R}^N , for example, by counting word frequencies
- so if you like you may assume that $X \subseteq \mathbb{R}^N$ where N is very large
- one way to do this:

- assume that

$$f(x) = \sum_{i=1}^n c_i K(x, x_i) \quad (4.1)$$

where $K : X \times X \rightarrow \mathbb{R}$ is some suitable *Mercer kernel*

- if $X \subseteq \mathbb{R}^N$ a common example is the Gaussian kernel

$$K(\mathbf{x}, \mathbf{y}) = e^{-\|\mathbf{x}-\mathbf{y}\|^2/2\sigma^2}$$

- since we already know the value of $f(x)$ for $x \in \{x_1, \dots, x_n\}$, we could in principle determine c_1, \dots, c_n by plugging x_1, \dots, x_n into (4.1) to get

$$f(x_1) = c_1 K(x_1, x_1) + \dots + c_n K(x_1, x_n)$$

$$f(x_2) = c_1 K(x_2, x_1) + \dots + c_n K(x_2, x_n)$$

...

$$f(x_n) = c_1 K(x_n, x_1) + \dots + c_n K(x_n, x_n)$$

or equivalently

$$\begin{bmatrix} K(x_1, x_1) & K(x_1, x_2) & \dots & K(x_1, x_n) \\ K(x_2, x_1) & \ddots & & \vdots \\ \vdots & & \ddots & \vdots \\ K(x_n, x_1) & \dots & \dots & K(x_n, x_n) \end{bmatrix} \begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_n \end{bmatrix} = \begin{bmatrix} f(x_1) \\ f(x_2) \\ \vdots \\ f(x_n) \end{bmatrix} \quad (4.2)$$

or

$$K\mathbf{c} = \mathbf{y}$$

- note that we know all values $K(x_i, x_j)$ of the matrix and also the right-hand side $f(x_i)$ as long as we have the training set $\{(x_i, f(x_i)) : i = 1, \dots, n\}$
- so we end up with a linear system like (2.6) again
- in principle this is very nice but in practice it rarely works since (4.2) is unlikely to have a solution
- so what we often need to do is to solve linear systems (2.6) approximately, i.e., $A\mathbf{x} \approx \mathbf{b}$ where ‘ \approx ’ is interpreted in some appropriate ways — we will look at some of these variants of (2.6) later
- the most common interpretation of $A\mathbf{x} \approx \mathbf{b}$ is the *least squares problem*

$$\min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|_2^2 = \min_{x_1, \dots, x_n \in \mathbb{R}} \sum_{i=1}^m \sum_{j=1}^n (a_{ij}x_j - b_i)^2 \quad (4.3)$$

- in the context of supervised learning, this is called *empirical risk minimization*, i.e., find c_1, \dots, c_n so that

$$\sum_{i=1}^n (y_i - f(x_i))^2 \quad (4.4)$$

is minimized

- but (4.4) is often *ill-posed* (no unique solution) and so a common strategy is to do Tikhonov regularization and minimize instead

$$\sum_{i=1}^n (y_i - f(x_i))^2 + \lambda \|f\|_K^2$$

where $\|\cdot\|_K$ is a special norm induced by the kernel $K(x, y)$ (if you must know, it is called the reproducing kernel Hilbert space or RKHS norm)

- as we will see later in this course, this leads to a problem of the form

$$(K + \lambda I)\mathbf{c} = \mathbf{y}$$

which is again a linear system except that we will need to find λ separately (we will see how to do this)

- now once we have c_1, \dots, c_n , given any x , we can find the value $f(x)$
- of course $f(x)$ would in general not be ± 1 but we can design a rule of the form

$$f(x) \begin{cases} < 0 & \Rightarrow x \text{ is spam} \\ > 0 & \Rightarrow x \text{ is not spam} \end{cases}$$

- so we have built a spam filter

5. SOLVING LINEAR SYSTEMS

- most of the course will focus on solving *linear systems* (2.6) and its variants like least squares, regularized least squares, total least squares, etc
- the fundamental problem is

$$A\mathbf{x} = \mathbf{b}$$

where we are given $A \in \mathbb{C}^{m \times n}$, $\mathbf{b} \in \mathbb{C}^m$ and we seek a solution $\mathbf{x} \in \mathbb{C}^n$

- often we will work over \mathbb{R} instead of \mathbb{C} but these would be only fields of interest
- some of the stuff we say in this course will be false over arbitrary fields (e.g. $\mathbb{F}_2 = \{0, 1\}$ with mod 2 arithmetic)
- three important numbers associated to a matrix A or a linear system $A\mathbf{x} = \mathbf{b}$:
 - m = number of rows = number of equations
 - n = number of columns = number of variables
 - $r = \text{rank}(A) = \dim(\text{im}(A)) = \dim(\text{colsp}(A)) = \dim(\text{rowsp}(A))$
- m, n, r tell us about existence and uniqueness of solution to $A\mathbf{x} = \mathbf{b}$
- terminologies
 - $m = n$: A is square matrix, $A\mathbf{x} = \mathbf{b}$ is a square system, i.e. number of variables equals number of equations

$$A = \begin{bmatrix} \times & \times & \times & \times \\ \times & \times & \times & \times \\ \times & \times & \times & \times \\ \times & \times & \times & \times \end{bmatrix}$$

- $m > n$: A is tall-and-thin matrix, $A\mathbf{x} = \mathbf{b}$ is an overdetermined system, i.e., more equations than variables

$$A = \begin{bmatrix} \times & \times \\ \times & \times \\ \times & \times \\ \times & \times \end{bmatrix}$$

- $m < n$: A is short-and-fat matrix, $A\mathbf{x} = \mathbf{b}$ is an underdetermined system, i.e., more variables than equations

$$A = \begin{bmatrix} \times & \times & \times & \times \\ \times & \times & \times & \times \end{bmatrix}$$

- if $r = \min\{m, n\}$, we say A is of *full rank*, otherwise we say A is *rank deficient*
- if $r = \min\{m, n\} = m$, we say A is of *full row rank*
- if $r = \min\{m, n\} = n$, we say A is of *full column rank*
- question: what is the big deal about solving linear systems $A\mathbf{x} = \mathbf{b}$? don't we know all about this already?

- answer: we only know how to solve idealized versions of the problem, but not in realistic situations
 - what if there are rounding errors in the coefficient matrix A or the right hand side \mathbf{b} or both
 - what if we want to solve it quicker than $O(n^3)$
 - what if m and n are large
 - what if we want to do things in parallel on multicore processors
 - what if we need to deal with a variant with constraints on the solution \mathbf{x} , or where $A\mathbf{x} = \mathbf{b}$ has no solution or no unique solution (as we saw in the machine learning example)
- linear systems are arguably the most widely solved problem in science and engineering
 - 70% of supercomputing time is spent on this
 - that's why solution of linear system is used to benchmark supercomputers (cf. <http://www.top500.org>)

6. TOP 10 ALGORITHMS OF THE 20TH CENTURY

- a broader motivation for this course and its sequel next quarter is that matrix computations are behind some of the most important algorithms
 - the three bold faced ones are algorithms in matrix computations
 - the four italics ones are algorithms are variants or extensions of algorithms in matrix computations
 - see <http://www.stat.uchicago.edu/~lekheng/courses/309/top10/>
- (1) Metropolis Algorithm for Monte Carlo
 - (2) *Simplex Method for Linear Programming*
 - (3) **Krylov Subspace Iteration Methods**
 - (4) **Decompositional Approach to Matrix Computations**
 - (5) Fortran Optimizing Compiler
 - (6) **QR Algorithm for Computing Eigenvalues**
 - (7) Quicksort Algorithm for Sorting
 - (8) *Fast Fourier Transform*
 - (9) *Integer Relation Detection*
 - (10) *Fast Multipole Method*

7. VARIANTS OF $A\mathbf{x} = \mathbf{b}$

- notations
 - $\mathbf{x} = [x_1, \dots, x_n]^T \in \mathbb{R}^n$

$$\|\mathbf{x}\|_2 := \sqrt{x_1^2 + \dots + x_n^2}$$

$$\|\mathbf{x}\|_1 := |x_1| + \dots + |x_n|$$

$$\|\mathbf{x}\|_\infty := \max\{|x_1|, \dots, |x_n|\}$$

$$\|\mathbf{x}\|_0 := \text{nnz}(\mathbf{x}) = \#\{i : x_i \neq 0\}$$

$$- A = [a_{ij}]_{i,j=1}^{m,n} \in \mathbb{R}^{m \times n}$$

$$\|A\|_F = \sqrt{\sum_{i,j=1}^{m,n} |a_{ij}|^2}$$

- we will discuss vector and matrix norms below
- note that $\|\cdot\|_0$ is not a norm

- (1) linear regression or least squares problem: know A exactly but \mathbf{b} is corrupted by error \mathbf{r} , i.e., $A\mathbf{x} = \mathbf{b} + \mathbf{r}$, and we want an \mathbf{x} that minimizes \mathbf{r} ,

$$\min\{\|\mathbf{r}\|_2^2 : A\mathbf{x} = \mathbf{b} + \mathbf{r}\} = \min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|_2^2 \quad (7.1)$$

Gauss–Markov theorem says that such an \mathbf{x} is the maximum likelihood estimator if the error \mathbf{r} is from a distribution that has zero mean and finite variance

- (2) error-in-variables regression or total least squares problem: A and \mathbf{b} are both corrupted by error E and \mathbf{r} , i.e., $(A + E)\mathbf{x} = \mathbf{b} + \mathbf{r}$, and we want an \mathbf{x} that minimizes both E and \mathbf{r} ,

$$\min\{\|E\|_F^2 + \|\mathbf{r}\|_2^2 : (A + E)\mathbf{x} = \mathbf{b} + \mathbf{r}\}$$

- (3) data least squares problem: A is corrupted by error E , i.e., $(A + E)\mathbf{x} = \mathbf{b}$, and we want an \mathbf{x} that minimizes E ,

$$\min\{\|E\|_F^2 : (A + E)\mathbf{x} = \mathbf{b}\}$$

- (4) minimum norm least squares: want the minimum length solution to (7.1),

$$\min\{\|\mathbf{x}\|_2^2 : \mathbf{x} \in \operatorname{argmin}\|A\mathbf{x} - \mathbf{b}\|_2^2\} = \min\{\|\mathbf{x}\|_2^2 : A^\top A\mathbf{x} = A^\top \mathbf{b}\} \quad (7.2)$$

the solution \mathbf{x}_* to (7.2) is unique and can in fact be used to define the Moore–Penrose pseudoinverse of A : $\mathbf{x}_* = A^\dagger \mathbf{b}$

- (5) robust regression: replace 2-norm by 1-norm (more generally, the Huber loss function) in (7.1),

$$\min\{\|\mathbf{r}\|_1 : A\mathbf{x} = \mathbf{b} + \mathbf{r}\} = \min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|_1$$

great for reducing sensitivity to outliers

- (6) ridge regression or regularized least squares

$$\min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|_2^2 + \|\Gamma \mathbf{x}\|_2^2$$

where $\Gamma \in \mathbb{R}^{p \times n}$ is some other matrix — most commonly $\Gamma = \lambda I$ or the finite-difference matrix

- (7) sparse or structured linear systems: sparse means A has a lot of zeroes (sufficiently many that it pays to take advantage of the fact), structured means that A can be defined with fewer than the usual number of mn parameters. An example of a data sparse matrix is a Toeplitz matrix

$$T = \begin{bmatrix} a_0 & a_1 & a_2 & & a_{n-1} \\ a_{-1} & a_0 & a_2 & \ddots & \\ a_{-2} & a_{-1} & \ddots & \ddots & a_2 \\ & \ddots & \ddots & \ddots & a_1 \\ a_{-n+1} & & a_{-2} & a_{-1} & a_0 \end{bmatrix} \in \mathbb{R}^{n \times n}$$

i.e., a_{ij} depends only on $|i - j|$, and T can be specified with just $2n - 1$ parameters $a_{-n+1}, \dots, a_{n+1} \in \mathbb{R}$; a Toeplitz $T\mathbf{x} = \mathbf{b}$ can be solved in $O(n \log^2 n)$ time as opposed to the usual $O(n^3)$ time for general linear systems

- (8) linear programming:

$$\min\{\mathbf{c}^\top \mathbf{x} : A\mathbf{x} \leq \mathbf{b}\}$$

note that $\mathbf{c}^\top \mathbf{x} = c_1 x_1 + \dots + c_n x_n$ is a linear function; this is very important in economics

- (9) quadratic programming: given $A \in \mathbb{R}^{n \times n}$, $B \in \mathbb{R}^{m \times n}$, $\mathbf{c} \in \mathbb{R}^n$, $\mathbf{d} \in \mathbb{R}^m$, want

$$\min \left\{ \frac{1}{2} \mathbf{x}^\top A \mathbf{x} - \mathbf{c}^\top \mathbf{x} : B\mathbf{x} = \mathbf{d} \right\}$$

this reduces to a linear system

$$\begin{bmatrix} A & B^T \\ B & 0 \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ \boldsymbol{\lambda} \end{bmatrix} = \begin{bmatrix} \mathbf{c} \\ \mathbf{d} \end{bmatrix}$$

- (10) basis pursuit: if we want the sparsest solution to an underdetermined linear system, we ought to solve

$$\min\{\|\mathbf{x}\|_0 : A\mathbf{x} = \mathbf{b}\}$$

but this is NP-hard and so we look at a convex relaxation

$$\min\{\|\mathbf{x}\|_1 : A\mathbf{x} = \mathbf{b}\}$$

which can in fact be reduced to a linear programming problem

8. NORMS

- a *norm* is a real-valued function on a vector space (over \mathbb{R} or \mathbb{C}), denoted $\|\cdot\| : V \rightarrow \mathbb{R}$ satisfying
 - (1) $\|\mathbf{x}\| \geq 0$ for all $\mathbf{x} \in V$
 - (2) $\|\mathbf{x}\| = 0$ if and only if $\mathbf{x} = \mathbf{0}$
 - (3) $\|\alpha\mathbf{x}\| = |\alpha|\|\mathbf{x}\|$ for all $\alpha \in \mathbb{C}$ and $\mathbf{x} \in V$
 - (4) $\|\mathbf{x} + \mathbf{y}\| \leq \|\mathbf{x}\| + \|\mathbf{y}\|$ for any $\mathbf{x}, \mathbf{y} \in V$
- we will be interested in two specific choices of V
 - $V = \mathbb{R}^n$ or \mathbb{C}^n
 - $V = \mathbb{R}^{m \times n}$ or $\mathbb{C}^{m \times n}$

9. VECTOR NORMS

- if $V = \mathbb{C}^n$ or $V = \mathbb{R}^n$, we call a norm on V a *vector norm*
- example: consider $\|\cdot\|_1 : \mathbb{C}^n \rightarrow \mathbb{R}$ defined by

$$\|\mathbf{x}\|_1 = \sum_{i=1}^n |x_i|$$

for $\mathbf{x} = [x_1, \dots, x_n]^T \in \mathbb{C}^n$ and where $|x|$ denotes the modulus/absolute value of $x \in \mathbb{C}$

– check that this is a norm:

- (1) clearly $\|\mathbf{x}\|_1 \geq 0$
- (2) the only way a sum nonnegative entries $\|\mathbf{x}\|_1 = 0$ is if all entries $|x_i| = 0$ and so $\mathbf{x} = [0, \dots, 0]^T = \mathbf{0}$
- (3) we have

$$\|\alpha\mathbf{x}\|_1 = \sum_{i=1}^n |\alpha x_i| = |\alpha| \sum_{i=1}^n |x_i| = |\alpha| \|\mathbf{x}\|_1$$

since complex modulus satisfies $|\alpha x| = |\alpha| |x|$

- (4) using the triangle inequality for complex numbers, we obtain

$$\|\mathbf{x} + \mathbf{y}\|_1 = \sum_{i=1}^n |x_i + y_i| \leq \sum_{i=1}^n |x_i| + |y_i| \leq \|\mathbf{x}\|_1 + \|\mathbf{y}\|_1$$

– therefore the function defines a norm, called the 1-norm or *Manhattan norm*

- example: more generally, for $p \geq 1$ (can be any real number, not necessarily an integer), we define the *p-norm* $\|\mathbf{x}\|_p$ by

$$\|\mathbf{x}\|_p = (|x_1|^p + \dots + |x_n|^p)^{1/p}$$

- most commonly used p -norms is the 2-norm or *Euclidean norm*:

$$\|\mathbf{x}\|_2 = \left(\sum_{i=1}^n |x_i|^2 \right)^{1/2}$$

- easy to see that for any p , we have

$$\left(\max_{i=1,\dots,n} |x_i|^p \right)^{1/p} \leq \|\mathbf{x}\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{1/p} \leq \left(n \max_{i=1,\dots,n} |x_i|^p \right)^{1/p}$$

- from which it follows that

$$\max_{i=1,\dots,n} |x_i| \leq \|\mathbf{x}\|_p \leq n^{1/p} \max_{i=1,\dots,n} |x_i|$$

- as $p \rightarrow \infty$, we obtain the *infinity norm*

$$\|\mathbf{x}\|_\infty = \lim_{p \rightarrow \infty} \|\mathbf{x}\|_p = \max_{i=1,\dots,n} |x_i|$$

which is also known as the *Chebyshev norm*

- easy to verify that p -norms for any $p \in [1, \infty]$ are indeed norms